

Bayesian Analysis of Randomized Controlled Trials

Julian Bautista, Alex Pavlakis, Advait Rajagopal

November 26, 2017

Abstract

Contents

1	Introduction	3
2	Bayesian Data Analysis	3
2.1	Model development	3
2.2	Model estimation	3
2.3	Model checking, comparison and expansion	4
3	Comparison of Bayesian Data Analysis to Other Methods	4
4	Reproducibility	4
5	Impact of Smartphone App on Eating Behavior	4
5.1	Experimental design	5
5.2	Exploratory data analysis	5
5.3	Model development	6
5.4	Model estimation	8
5.5	Posterior predictive checking	8
5.6	Results	8
6	Conclusion	10
7	Appendix	10

1 Introduction

Set up the paper. This is me learning GitHub.

2 Bayesian Data Analysis

Bayesian data analysis (BDA) and inference is the process of developing and fitting a probability model to data. The result is learning the probability distribution of the parameters of the model and being able to evaluate the fit of the model to the data as well being able to make predictions for new observations [Gelman et. al 2013; Gelman and Hill 2007].

There are three main steps to BDA and these are listed below and explained in detail in section 2.1, 2.2 and 2.3 respectively.

1. Set up a probability model.
2. Calculate the distribution of the model parameters, given the observed data. This is called the *posterior distribution* of the parameters.
3. Check the fit of the model, whether the conclusions are sensible and how sensitive they are to modeling assumptions. Expand or alter the model if needed and repeat the steps.

2.1 Model development

Developing a probability model means setting up a joint probability distribution in a way that it takes into account all observed data and unobserved quantities. It should include all knowledge of the experiment or data collection process and should be logically consistent with scientific nature of the problem at hand. This involves a *prior distribution* and a *likelihood function*.

Bayesian analysis requires the assignment of a prior distribution that can restrict the set of values that a coefficient or parameter can take. This prior distribution also allows us to make our assumptions about the experiment and actual nature of the problem explicit. For example, if we were dealing with a regression prediction problem, we could assign a completely “noninformative” prior to our coefficient which is equivalent to saying that the coefficient is a draw from a uniform distribution on the whole real line and moreover that all those values are equally likely. Incidentally this boils down to a maximum likelihood estimate of the coefficient. But statisticians are rarely in a position where they know absolutely nothing about the nature of their model parameters. Perhaps they know it needs to be positive, or if what is being modeled is a probability then it needs to be between zero and one. Either way, this knowledge about the parameter should be made explicit and incorporated into the problem using an “informative” prior distribution. In section 3 we discuss some advantages and choices of prior distributions. The likelihood function is the probability distribution of the observed data conditional on the parameters. Again, in a regression problem context, the likelihood function is a Gaussian distribution.

The Bayesian approach combines data and prior information, which coupled with the correct model yield more precise inferences than we could get from either source on its own.

2.2 Model estimation

The goal of model estimation is to calculate the *posterior distribution* of the parameters. This is the conditional distribution of the parameter given the observed data. The posterior distribution is obtained by simply multiplying the prior and likelihood to get joint

probabilities and then normalizing to get posterior probabilities. Calculating the posterior analytically is often difficult and leads to intractable integrals with no closed-form solution. So the standard practice is to use MCMC sampling methods to approximate the posterior up to a normalizing constant and sample from it. There are some other approaches to calculating the posterior distribution, but those are beyond the scope of this paper. Using these posterior distributions we can now predict and simulate replicated data from our model, which leads us to section 2.3.

2.3 Model checking, comparison and expansion

We begin evaluating the performance of our model by ensuring that it converges on robust parameters and that the results make sense given what we know about the observed data and the problem itself. We also need to see how sensitive our results are to the modeling assumptions and the priors. Sometimes using stronger priors can help with convergence and more stable results. We use our posterior distributions of the parameters to check our model by carrying out *posterior predictive checks*. This could be graphical checks where we simulate “fake” data from our model and compare it graphically to the raw data. We could also use certain test statistics and compare the values of those statistics across models and see how close they are to the ground truth obtained from the observed data itself. Based on the fit of the model, we can either expand the model to include more parameters, predictors or change the model and reevaluate performance. We are not looking to explicitly choose the “right” model but are interested to learn where each of our approaches and models might be lacking and make these shortcomings explicit.

3 Comparison of Bayesian Data Analysis to Other Methods

Write about the advantages of BDA over other formats. How are they similar? How are they different?

4 Reproducibility

Describe the philosophy and importance of reproducibility.

Describe the steps of making research reproducible (e.g. sharing code, sharing data, explicitly specifying model, etc.)

5 Impact of Smartphone App on Eating Behavior

Hildebrandt et al (2017) conducted an experiment to test whether the Noom Monitor, a smartphone application, could augment the effect of in-person therapeutic treatment on binge eating behavior. The treatment, known as *guided self-help treatments based on cognitive-behavior therapy* (CBT-GSH), had been shown in previous research to reduce binge eating behavior by 10-50%. The Noom Monitor application was designed to facilitate CBT-GSH. For this example, we consider two research questions from the experiment:

1. Is CBT-GSH more effective at reducing binge eating behavior when facilitated by the Noom Monitor?
2. Does the effect of the Noom Monitor vary over time?




Figure 1: *Plots display OBEs over time for each individual in the CBT-GSH groups (left panel) and CBT-GSH + Noom group (right panel).*

5.1 Experimental design

66 men and women with Bulimia Nervosa (BN) or Binge Eating Disorder (BED) were randomly assigned into two treatment conditions: CBT-GSH (N= 33) or CBT-GSH + Noom (N=33). Therapy lasted for 12 weeks. Assessments were conducted at weeks 0, 4, 8, 12, 24, and 36. The primary outcome was Objective Bulimic Episodes (OBE).

5.2 Exploratory data analysis

Figure X displays OBEs per week for each individual in both treatment conditions. A few aspects of the data immediately stand out, which suggest that any model should account for individual-level effects and time-level effects, and should let treatment effects vary over time.

- The number of OBEs decreases over the course of the treatment for almost all subjects
- The biggest decreases in OBEs appear to occur in the early stages of treatment
- The primary sources of variation in OBE appear to be *between people* and *over time*.




Figure 2: *Histograms display the distribution of OBEs in each condition in each week.*

Figure Y displays the distribution of histograms in each condition in each week. The distributions appear to condense around zero for both conditions over time, and the distributions in the CBT-GSH condition appear to have longer tails than those in the CBB-GSH+Noom condition.

5.3 Model development

We analyze RCTs by modeling the outcome of interest (in this case OBE) as a function of the treatment and all available pre-treatment covariates. The coefficients associated with the treatment reveal the average treatment effects. Inclusion of all available pre-treatment covariates accounts for variation in the outcome variable, decreasing uncertainty around treatment effects and providing the model with more predictive power. We conduct *intent-to-treat* analysis, meaning that our inferences will be based on initial treatment assignment, and will not account for mid-experiment dropouts.

The outcome variable is restricted to be nonnegative integers, so we fit a Poisson regression model, with hierarchies on individuals, time periods, and treatment effects. For each individual in each time period, the number of OBEs follows a Poisson distribution, with a mean dependent on the characteristics of the individual and the time period.

$$OBE_{i,t} \sim \text{Poisson}(\lambda_{i,t}) \quad (1)$$

$$\lambda_{i,t} = \exp(\alpha_i + \beta_t + \gamma_t T_i + X_i \theta) \quad (2)$$

$$i = 1, \dots, 66 \quad (3)$$

$$t = 0, 4, 8, 12, 24, 36 \quad (4)$$

$$T_i = \begin{cases} 0, & \text{if } CBT - GSH \\ 1, & \text{if } CBT - GSH + Noom \end{cases} \quad (5)$$

α is an individual-specific intercept, β is a time-specific intercept, γ is a time-specific treatment effect, T is a treatment indicator, X is a matrix of individual level covariates (age, sex, race, etc), and θ is a vector of individual level effects.

We believe that individual-level intercepts are simultaneously unique to the individual and common to the population; that is, each individual has their own baseline predilection to engage in eating disorder behavior, but their baseline predilections are not drastically different from each other. We operationalize this concept by modeling all individual-level intercepts as coming from a common distribution, with *hyperparameters* μ_α and τ_α .

$$\alpha_i \sim \text{Normal}(\mu_\alpha, \tau_\alpha) \quad \forall i \in 1, \dots, 66 \quad (6)$$

Similarly, we believe that time-specific treatment effects may be unique to each period but similar over time. We operationalize this concept by modeling all time-specific treatment effects γ as coming from a common distribution, with *hyperparameters* μ_γ and τ_γ .

$$\gamma_t \sim \text{Normal}(\mu_\gamma, \tau_\gamma) \quad \forall t \in 0, 4, 8, 12, 24, 36 \quad (7)$$

μ_γ is the *grand mean*, the overall treatment effect; τ_γ is the variation in treatment effects over time; and each individual γ_t is a time-period specific treatment effect. This approach has a natural smoothing effect: any extreme estimates of γ_t will be partially-pooled back toward the grand mean μ_γ .

To complete the model we must assign prior distributions to parameters (such as θ) and hyperparameters (such as μ). These prior distributions will serve three main purposes:

1. Transparency: priors allow us to be clear about our modeling assumptions.
2. Convergence: priors help the model "know where to look" and converge.
3. Inclusion of Additional Information: We usually have more information about an experiment than is represented by the data. Priors allow us to include that information explicitly in our model.

For this model, we assign the following prior and hyperprior distributions:

$$\mu_{\alpha} \sim \text{Normal}(5, 5) \quad (8)$$

$$\tau_{\alpha} \sim \text{Cauchy}^+(0, 30) \quad (9)$$

$$\mu_{\gamma} \sim \text{Normal}(0, 5) \quad (10)$$

$$\tau_{\gamma} \sim \text{Cauchy}^+(0, 30) \quad (11)$$

$$\theta \sim \text{Normal}(0, 1) \quad (12)$$

$$(13)$$

The normal distributions around the individual and treatment effects allow us to guide the model to the appropriate range of parameter values, but with wide enough variance (5 in each case) to let the model find its own way in that range. Half cauchy priors on the variance parameters are weakly informative, with much of their mass around zero but gentle slopes in their tails, which have been shown to be effective prior distributions for variance parameters (Gelman, 2006).

5.4 Model estimation

We estimate this model with *Hamilton Monte-Carlo* in Stan. Model code is appended to this document. We find that the model converges with four chains of 2000 iterations each (see table M).

5.5 Posterior predictive checking

Before using our model to make inferences about time-specific treatment effects, we check its fit by comparing model-simulated OBE to data OBE. If model simulations do not track the data well, we may want to revisit our model’s assumptions before trusting its inferences. If the model’s simulations recover patterns in the data, we are more inclined to trust its inferences.

Figure Q displays OBEs in each period for each individual in each treatment group, from raw data (upper plots) and model simulations (lower plots). Black lines display means for each. This suggests that the model is broadly able to pick up on the key variables that determine OBE over time for the duration of this experiment.

Another way to check the fit of the model is by comparing simulated data directly against the raw data. Figure P shows this for both treatment conditions. Simulated data for the Noom condition appears to better track the raw data than simulated data for the no Noom condition. This is unsurprising, since the no Noom condition tended to have more outliers, which we would not expect (or want) our model to pick up perfectly from such a small sample.

We conclude our model evaluations by mapping modeled density curves for each condition in each time period over the histograms in figure Y. Figure BB shows that our model is able to broadly pick up on the patterns in the data over time and between treatment conditions.

5.6 Results

Model results are displayed in table M. Results suggest that using the Noom Monitor smartphone application during CBT-GSH may slightly decrease OBEs. There some evidence that




Figure 3: *Upper plots show OBEs in each period for each individual in both treatment group, with black lines representing means. Lower plots show modeled OBEs for each individual in each treatment group, with black lines representing modeled means. The model appears to be able to recover OBEs over time fairly well for both treatment conditions.*




Figure 4: *Plots display simulated OBEs to raw data for both treatment conditions. Model simulations appear to match the raw data well, particularly for the Noom condition, which had fewer outliers.*

the treatment effect varies over time, with the Noom effect being slightly more pronounced during the later stages of therapy.

Figure R displays modeled OBE for both treatment groups (upper plot) and smoothed treatment effects (lower plot). In each measurement period, simulated OBE are higher for the No Noom condition than for the Noom condition, with some of the difference likely attributable to use of the Noom Monitor smartphone app.

6 Conclusion

So awesome

7 Appendix

All the math goes here

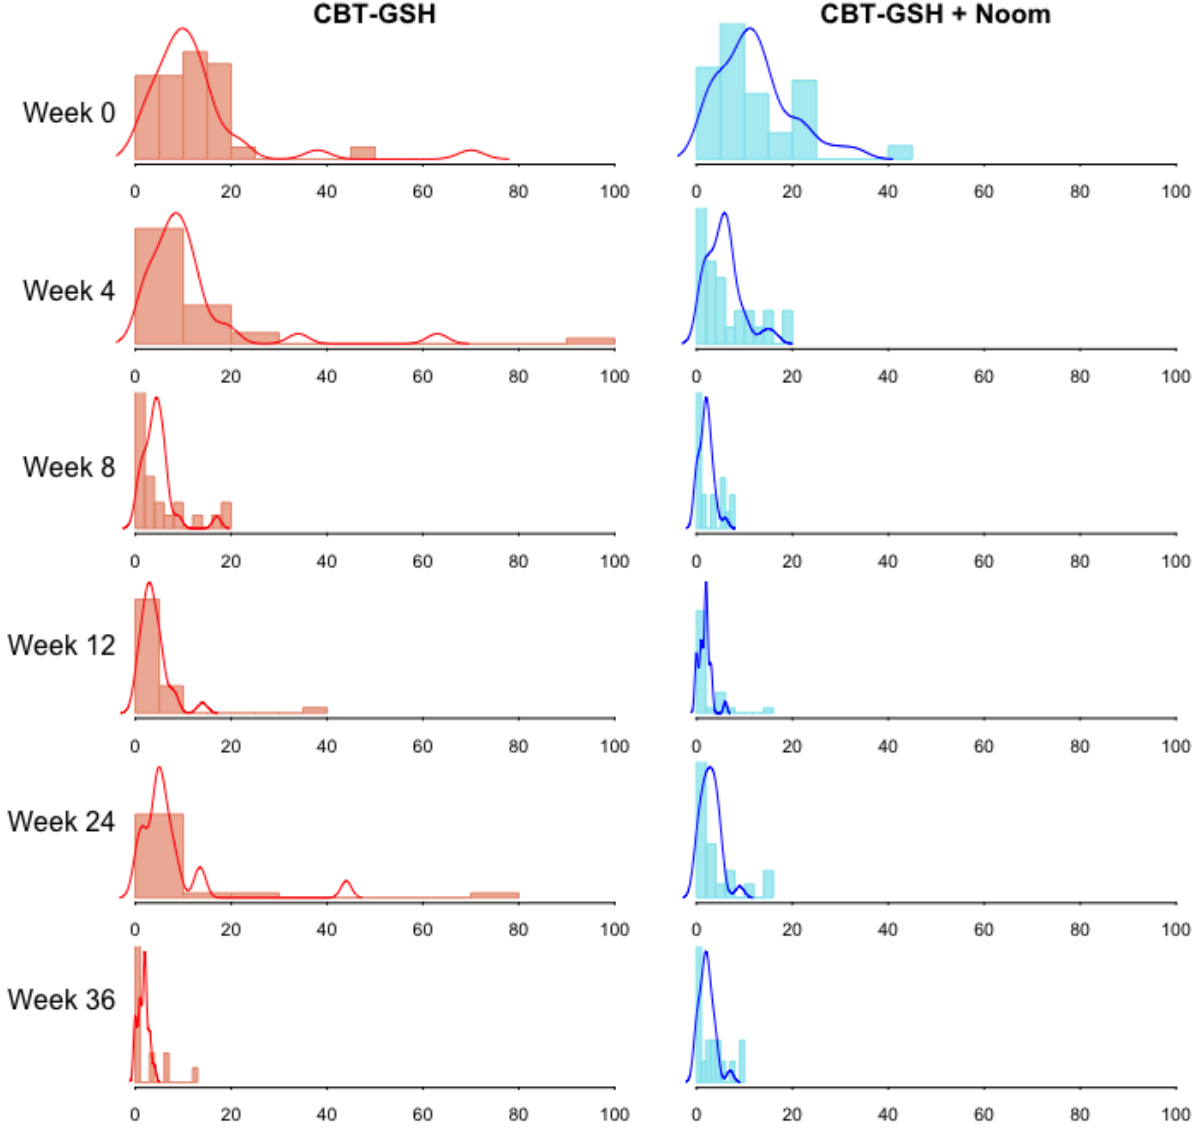


Figure 5: Histograms display the distribution of OBEs in each condition in each week.

	mean	25%	50%	75%
γ_0	0.18	-0.45	0.15	0.78
γ_4	-0.43	-1.05	-0.46	0.16
γ_8	-0.70	-1.33	-0.71	-0.10
γ_{12}	-0.65	-1.28	-0.68	-0.04
γ_{24}	-0.72	-1.34	-0.75	-0.11
γ_{36}	0.21	-0.42	0.19	0.82
μ_γ	-0.34	-0.98	-0.36	0.26
τ_γ	0.64	0.43	0.56	0.77

Table 1: Table displays model results for Noom effects in all six time periods and grand mean and variance parameters.




Figure 6: Upper plot displays modeled OBEs in each time period for the Noom (blue) and no Noom (orange) conditions with 95% intervals. Lower plot displays modeled treatment effects in each period, with 50% intervals.